FOREX TRADING

The Forex Trading Playbook

Building Your AI-Powered Support System

Volume 11 — AI Agents for the Forex Trader

Saeed Mahmoud Jajah

Edition 2026

Contents

Read in order on your first pass — each chapter builds on the last — then keep this as a reference.

Introduction

AI Agents for the Forex Trader — Building Your AI-Powered Support System

Every trader now has access to something the legends in **Vol 2** never had: a tireless, knowledgeable assistant that can research, organize, analyze, quiz, and hold you accountable — on demand, at near-zero cost. Used well, **AI agents** can compress the years it takes to become consistent by removing friction from the disciplines that actually matter: preparation, journaling, review, and research. Used badly, they become just another expensive way to lose money chasing predictions.

This volume is the honest, practical guide to building a **team of AI agents** around your trading — not to *trade for you*, but to make *you* a sharper, more disciplined, more consistent trader. It maps a specific set of agents onto the daily routine of **Vol 10**, gives you ready-to-paste prompts for each, and draws a hard line around what you must **never** automate.

> ⚠️ **The one rule that governs this entire volume.** AI agents **cannot predict the market**, and any agent (or vendor) that claims to is the modern face of the scams in **Vol 1**. Your edge is not the AI — your edge is *discipline, risk control, and process*. AI agents are a force-multiplier on that process. They support your judgment; they never replace it. Hand an AI your decision-making and you've automated away the only thing that was ever going to make you money.

This is **Volume 11**. It turns the analysis of the whole series — the routine of **Vol 10**, the journaling system of **Vol 4**, the macro process of **Vol 6**, and the trading plan of **Vol 7** — into an AI-augmented operating system you actually run.

Your AI agent team in one diagram

```
┌──────────────────────────────────────────────────────────┐
│                 YOU — the decision-maker                  │
│       (judgment, discipline, risk, final say on every trade)      │
└──────────────────────────────────────────────────────────┘
       ▲             ▲             ▲            ▲          ▲
PRE-MARKET      SESSION       POST-MARKET    OFF-HOURS   ONGOING
┌─────────┐  ┌───────────┐  ┌───────────┐  ┌────────┐  ┌────────┐
│ BRIEFING│  │ GATEKEEPER│  │ JOURNAL & │  │ REVIEW │  │ MACRO  │
│  agent  │  │   agent   │  │ ANALYTICS │  │COACHING│  │RESEARCH│
│ (Ch. 3) │  │  (Ch. 4)  │  │  (Ch. 5)  │  │ (Ch. 7)│  │ (Ch. 6)│
└─────────┘  └───────────┘  └───────────┘  └────────┘  └────────┘
prep the day   vet each trade   log & measure   find lessons  digest news
    │             │               │             │          │
    └─────────────┴──── all feed ─┴─────────────┴──────────┘
                 your decisions — never make them
```

> 🔑 **AI does the work; you make the calls.** Every agent in this volume gathers, organizes, checks, and reflects — turning hours of friction into minutes — so that *you* arrive at each decision better prepared and more disciplined. The trader still trades. The AI just makes sure the trader shows up prepared, follows the rules, and learns from every trade.

CHAPTER 1

The AI-Augmented Trader

Before you build a single agent, you need the right mental model — because the difference between AI *helping* your trading and AI *destroying* your account is entirely a question of how you frame it. This chapter draws that line clearly and honestly. Get this chapter wrong and every clever prompt in the rest of the volume becomes a faster way to lose money. Get it right and AI becomes the quiet force-multiplier behind your discipline.

> 🔑 **AI is an augmenter of process, not a generator of edge.** It will not tell you where the market is going. It *will* help you prepare better, follow your rules more reliably, journal everything, spot your own patterns, and research faster than you ever could alone. The edge was always going to come from *you* — your discipline and risk control. AI just removes the friction that makes discipline hard.

1.1 What AI agents genuinely do well for trading

Large language models (LLMs) and the agents built on them are extraordinary at a specific class of tasks — and trading is full of them:

```
AI IS GENUINELY EXCELLENT AT:
  • SUMMARIZING — turning a 12-page central-bank statement into 5 key points (Vol 6)
  • STRUCTURING — taking your messy trade notes and formatting them into a clean journal
  • CHECKING — comparing a proposed trade against a fixed checklist, tirelessly, every time
  • CALCULATING — position size, R-multiples, expectancy, KPIs from your trade data (Vol 4)
  • RESEARCHING — gathering and organizing information on a currency, event, or theme
  • EXPLAINING — answering "what does an inverted yield curve mean for the dollar?"
  • QUIZZING — drilling you on your own playbooks and trade reasoning
  • REMEMBERING — holding your rules, your plan, and your history and reflecting them back
  • REDUCING FRICTION — making the disciplined behavior the easy, fast, default behavior
```

Notice what unites these: they are all **process-support** tasks. None of them is "predict the market." That is not an accident — it's the whole thesis.

1.2 What AI agents cannot do (and never will)

Equally important — arguably *more* important — is an honest list of the limits:

```
AI CANNOT (no matter what anyone sells you):
  • PREDICT price — markets are not a text-completion problem; no model has an edge here
  • GUARANTEE anything — "risk-free," "90% win rate" AI = the Vol 1 scams in a new costume
  • REPLACE your judgment — it has no skin in the game and no accountability for your losses
  • KNOW what it doesn't know — it will state false things confidently ("hallucinate")
  • ACCESS live truth by default — a plain LLM has a knowledge cutoff and no live prices
  • FEEL the consequences — only YOU lose the money, so only YOU can own the decision
```

> ⚠️ **The hallucination problem is non-negotiable for traders.** An LLM can invent a data release, misquote an interest rate, or fabricate a "fact" with total confidence. **Never act on an AI's factual claim without verifying**

it at the source (the actual economic calendar, the actual chart, the actual central-bank page). Treat every agent output as a *draft for your review*, not a truth.

1.3 The two mindsets — and why one of them blows up accounts

There are two completely different ways to bring AI into trading. They lead to opposite outcomes:

	❌ The "Oracle" mindset (blows up)	✅ The "Co-pilot" mindset (compounds)
Sees AI as	A predictor / signal source	A research & discipline assistant
Asks it	"Should I buy EUR/USD now?"	"Summarize today's events; check this trade vs my rules"
Decision	Outsourced to the machine	Owned by the trader
Edge from	Hoping the AI is right	Better prep, tighter discipline, faster learning
When it's wrong	You lose, blame the AI, repeat	You caught it, because you verified
Outcome	Dependency, blown account	A sharper, more consistent trader

🔑 **The Co-pilot flies; the pilot commands.** A co-pilot handles the checklists, the radio, the navigation math, and the monitoring — freeing the pilot to fly and decide. It never takes the controls in a crisis. Build every agent in this volume as a co-pilot, and you'll be fine. Build one as an oracle and it's only a matter of time.

1.4 Why AI is a genuine advantage *if* used right

None of the above means AI is useless — quite the opposite. Used as a co-pilot, it is one of the biggest advantages a modern retail trader has ever had:

- **It compresses the learning curve.** The grind of becoming consistent (Vol 10) is mostly friction: preparing, journaling, reviewing, researching. AI strips the friction so you actually *do* these things daily instead of skipping them.
- **It democratizes the back office.** What used to require an analyst team — daily macro briefings, performance analytics, research digests — a solo trader can now approximate with well-built agents.
- **It enforces consistency.** A checklist you might rush, the Gatekeeper agent (Ch. 4) runs identically every single time, immune to FOMO and fatigue.
- **It's a patient teacher.** It will explain, re-explain, quiz, and review your reasoning without judgment, on your schedule.

The trader who pairs **genuine discipline** with a **well-built AI support system** has a real, durable advantage over the trader who has neither — or who has outsourced their judgment to a black box.

1.5 The prime directive for this volume

Everything that follows rests on one sentence. Tape it to your monitor:

```
┌──────────────────────────────────────────────────────────┐
│  AI agents support my PROCESS. I own every DECISION.        │
```

```
│  They prepare, check, log, research, and reflect.                    │
│  I judge, I risk, I decide, I am accountable.                        │
│  I verify their facts. I never trade on their prediction.            │
└──────────────────────────────────────────────────────────────────────┘
```

With that frame locked in, the rest of this volume is safe — and genuinely powerful. We'll now design the *architecture* of your agent team (Ch. 2), then build each agent one by one.

1.6 Chapter summary

```
• AI augments your PROCESS; it does not generate market EDGE. Your edge is still
  discipline + risk control. AI removes the friction that makes discipline hard.
• AI excels at process-support: summarizing, structuring, checking, calculating,
  researching, explaining, quizzing, remembering, reducing friction.
• AI CANNOT predict price, guarantee outcomes, replace your judgment, or know when
  it's wrong. It hallucinates — VERIFY every factual claim at the source.
• Two mindsets: the "Oracle" (asks AI what to do → dependency → blow-up) vs the
  "Co-pilot" (AI assists, trader decides → discipline → compounding). Be the Co-pilot.
• Used right, AI is a real advantage: it compresses the learning curve, gives a solo
  trader a "back office," enforces consistency, and teaches patiently.
• PRIME DIRECTIVE: AI supports my process; I own every decision. I verify its facts;
  I never trade on its prediction.
```

CHAPTER 2

The Agent Architecture

You don't want *one* do-everything AI; you want a **small team of specialized agents**, each with a narrow job, mapped onto the moments of your trading day. This chapter gives you the architecture — what an "agent" actually is, which agents to build, in what order, and how they connect into the routine of **Vol 10**.

> 🔑 **Specialized beats general.** A single agent told to "help me trade" gives vague, mediocre output. Five agents each with one sharp job — brief me, vet this trade, log this, research this, review my week — each give excellent output, because each has a clear role, a fixed format, and a focused context. Build narrow specialists, not one generalist.

2.1 What "an agent" actually means here

Strip away the hype and an AI "agent," for your purposes, is just three things:

```
AN AGENT =  ROLE + INSTRUCTIONS + (optionally) TOOLS & MEMORY

  ROLE          a clear identity and single job ("You are my pre-market briefing analyst")
  INSTRUCTIONS  exactly what to do, what format to output, and what NOT to do (its rules)
  TOOLS         optional access to live data (web, calendar, your files, a price API)
  MEMORY        optional persistence — your rules, your plan, your trade history
```

In practice an agent can be as simple as **a saved prompt you paste into a chat**, or as advanced as **a script that calls an LLM API, pulls live data, and writes to your journal file**. Start simple. A well-written prompt you actually use beats a sophisticated system you never finish building (Ch. 8).

```
MATURITY LADDER (start at the top; climb only as needed):
  LEVEL 0  A saved prompt you paste into a chat app each day           ← start here
  LEVEL 1  A "custom" reusable assistant with a fixed system prompt
  LEVEL 2  + tools: it can search the web / read your files / a calendar
  LEVEL 3  + memory & data: it reads your journal, knows your plan
  LEVEL 4  + orchestration: a script chains agents and writes outputs    ← advanced
```

2.2 The support-layer model

The core architectural idea: AI sits as a **support layer wrapped around you**, the decision-maker. It never sits *between* you and the market.

```
          ❌ WRONG (oracle)                    ✅ RIGHT (co-pilot)

     MARKET → AI → YOU → order          MARKET → YOU → order
                                                  ▲
       AI decides, you click              AI support layer feeds YOU:
                                            briefing · checks · logs ·
                                            research · review
```

Every agent's output flows *to you* to improve your decision. No agent's output flows *to the market* without you in the loop. This single rule keeps the whole system safe (Ch. 1).

2.3 The five core agents (and how they map to your day)

Map the agents directly onto the **Vol 10** routine — that's what makes them stick, because they attach to cues you already have:

Agent	When (routine cue)	Its one job	Chapter	Feeds from
Pre-Market Briefing	Morning prep	Assemble a structured daily briefing	Ch. 3	Vol 3, 6
Trade-Plan Gatekeeper	Before each entry	Vet the trade against your rules	Ch. 4	Vol 4, 7
Journal & Analytics	Post-market close	Log trades & compute KPIs	Ch. 5	Vol 4
Review & Coaching	Weekly / monthly	Find patterns, teach, plan	Ch. 7	Vol 4, 7, 10

🔑 **Anchor each agent to an existing routine cue and it becomes automatic.** "Coffee → run the Briefing agent." "Before I click buy → run the Gatekeeper." "Session close → feed the Journal agent." "Sunday → run the Review agent." The agents don't add a new routine; they plug into the one you already have from Vol 10.

2.4 The build order — Tier by tier

Don't build all five at once (that's the Vol 10 over-reach mistake). Build in tiers, mastering each before adding the next:

```
TIER 1 — build first (highest leverage, lowest risk):
  1. JOURNAL & ANALYTICS  — the single highest-ROI agent; start logging immediately
  2. PRE-MARKET BRIEFING  — removes the biggest daily friction (prep)
  3. TRADE-PLAN GATEKEEPER — your discipline enforcer at the point of decision

TIER 2 — add as you mature:
  4. MACRO & NEWS RESEARCH — depth on fundamentals (pairs well with Vol 6)
  5. REVIEW & COACHING     — runs on your now-rich journal data

TIER 3 — optional / advanced (Ch. 8):
  • Backtest & strategy-analysis helper · Accountability/habit checker
  • Orchestration that chains the above
```

⚠️ **Resist the urge to automate the decision as you climb tiers.** The more capable your system gets, the more tempting it is to let it "just take the obvious trades." Don't. Capability is not permission. The prime directive (Ch. 1) holds at every tier: agents support, you decide.

2.5 Three things every agent needs from you

Whatever the agent, it can only be as good as what you give it. Every agent in this volume depends on three inputs you must supply:

1. **Your rules & plan (context).** The Gatekeeper can't check a trade against rules it doesn't have. Feed each agent the relevant parts of your **Vol 7** trading plan — your strategy, risk limits, and checklist. *Garbage-in still equals garbage-out.*
2. **A fixed output format.** Tell the agent *exactly* how to respond (a template, a table, a checklist). Consistent format = output you can actually use day after day and compare over time.
3. **Verification by you.** Every factual claim (a data time, a price level, a rate) is a *draft* until you confirm it at the source (Ch. 1). The agent drafts; you ratify.

2.6 A note on tools, data, and privacy

We cover the build details in **Ch. 8**, but architecturally, know the trade-offs now:

- **A plain chat LLM** (Level 0–1) has no live data — great for structuring, checking, explaining, and reviewing *information you paste in*. It does **not** know today's price or news.
- **An agent with tools** (Level 2+) can fetch live news, an economic calendar, or prices — far more powerful, but you must verify sources and watch costs.
- **Privacy:** your journal and plan are sensitive. Prefer providers/settings that don't train on your data; never paste broker credentials, account numbers, or anything you wouldn't want stored (Ch. 8).

2.7 Chapter summary

```
• An AGENT = ROLE + INSTRUCTIONS + (optional) TOOLS & MEMORY. It can be as simple as a
  saved prompt or as rich as a script with live data. Start simple (Level 0).
• SUPPORT-LAYER model: AI wraps around YOU and feeds your decisions; it never sits
  between you and the market. Output flows to you, never to the order, without you.
• FIVE core agents map to your Vol 10 routine: Briefing (prep), Gatekeeper (entry),
  Journal (close), Research (off-hours), Review (weekly/monthly).
• Anchor each agent to an existing routine CUE so it becomes automatic.
• BUILD ORDER: Tier 1 = Journal, Briefing, Gatekeeper. Tier 2 = Research, Review.
  Tier 3 = advanced/orchestration. Master each tier before the next.
• Every agent needs three things from you: your RULES/PLAN as context, a FIXED output
  FORMAT, and your VERIFICATION of its facts. Capability is never permission to decide.
```

CHAPTER 3

The Pre-Market Briefing Agent

The pre-market routine (**Vol 10, Ch. 3**) is where your edge is *built* — before any trade. But it's also friction-heavy: calendar, overnight news, risk regime, levels, bias. On a tired morning it's tempting to skip. The **Pre-Market Briefing Agent** removes that friction by assembling a structured briefing so you start every day prepared, not staring at a blank screen.

> 🔑 **The briefing agent's job is preparation, not prediction.** It does *not* tell you what will happen or what to trade. It *gathers and organizes* what you need to form your own bias and plan — turning 45 minutes of scattered prep into a 10-minute review of a clean, consistent briefing. You still form the bias. You still write the plan.

3.1 What the briefing agent produces

A single, consistent **daily briefing** in a fixed format you review each morning:

```
┌─ DAILY BRIEFING — [date] ─────────────────────────────────────────┐
│ 1. CALENDAR    Today's high-impact events for my pairs + exact times    │
│ 2. OVERNIGHT   What moved while I slept (Asian session) + why, briefly  │
│ 3. RISK REGIME Risk-on / risk-off read: DXY, yields, equities, VIX, gold │
│ 4. KEY LEVELS  Pre-marked support/resistance on my watchlist pairs      │
│ 5. THEMES      The current macro narrative in play (1–3 bullets, Vol 6) │
│ 6. WATCH       Pairs/events most worth attention today                  │
│ 7. PROMPTS     Questions for ME to answer to form my own bias & plan    │
└──────────────────────────────────────────────────────────────────────────┘
```

Note section 7: a *good* briefing agent ends by handing the decision back to you — "Based on the above, what's your bias on EUR/USD and what would invalidate it?" It prepares you to think; it doesn't think for you.

3.2 Two ways to build it

```
LEVEL 0–1 (no live data) — you paste, it structures:
  You gather the raw inputs (calendar screenshot, overnight notes, level list)
  and paste them in. The agent organizes them into the briefing format above,
  flags conflicts, and asks you the bias-forming questions.
  → Simple, private, reliable. The agent never invents data — YOU supplied it.

LEVEL 2+ (with tools) — it fetches, you verify:
  The agent has web/calendar/price tools and assembles the briefing itself.
  → Far less effort, but you MUST verify every event time and number at the
    source. Treat it as a draft analyst report, not gospel.
```

> ⚠️ **The data-accuracy line.** A high-impact event at the wrong time can wreck your day (you trade into an NFP you thought was an hour away). If the agent fetches the calendar, **cross-check every high-impact event against the official source** before you rely on it. The fastest place for hallucination to hurt you is a wrong event time.

3.3 The core prompt (starter)

A Level 0–1 version you can use today — full versions are in the **Prompt Pack**:

```
ROLE: You are my pre-market briefing analyst. You prepare me to trade; you never
predict the market or tell me what to trade.

MY CONTEXT:
  • Pairs I trade: [e.g. EUR/USD, GBP/USD, USD/JPY]
  • My session: [e.g. London–NY overlap]
  • My style/bias inputs: [scorecard / trend method from Vol 3 & 6]

INPUTS I'll paste: today's economic calendar, overnight price action notes,
current DXY/yields/VIX read, and my marked key levels.

DO:
  1. Organize everything into my fixed BRIEFING format (7 sections).
  2. Flag any conflicts or notable risk events explicitly.
  3. Keep it concise and factual — only what I gave you; do not invent data.
  4. End with 3–4 sharp questions that help ME form my own bias and day plan.

DON'T: predict direction, give buy/sell calls, or state facts I didn't provide.
```

3.4 How it plugs into your routine

```
ROUTINE CUE: "coffee → run the briefing agent"  (Vol 10, Ch. 3 & 8)

  07:00  Wake, hydrate, move (Vol 10, Ch. 2)
  08:30  Gather raw inputs → run Briefing agent → review the briefing (10 min)
  08:45  Using section 7's questions, write MY OWN bias + day plan (Vol 10, Ch. 3)
  09:30  At the screen, prepared — hand the plan to the Gatekeeper (Ch. 4)
```

The agent compresses the *gathering*; you keep the *thinking*. You still produce a written bias and an if-then plan in your own words — that act of writing is the prep that builds the edge (Vol 10).

3.5 Pitfalls to avoid

```
✗ Letting the briefing REPLACE your analysis — it's an input to your thinking, not a substitute
✗ Trusting fetched event times/numbers without verifying at the source
✗ Asking it "so what should I trade?" — that's the oracle trap (Ch. 1); it prepares, you decide
✗ A bloated 3-page briefing you won't read — keep the format tight and skimmable
✗ Skipping the section-7 questions — that's where YOU do the actual prep work
```

3.6 Chapter summary

```
• The Briefing agent automates the GATHERING of pre-market prep (Vol 10, Ch. 3) into a
  fixed, skimmable daily briefing — calendar, overnight, regime, levels, themes, watch.
• It PREPARES you; it never predicts or gives buy/sell calls. It ends by asking YOU the
  questions that form your own bias and plan.
• Build it Level 0–1 (you paste inputs, it structures) for privacy & reliability, or
  Level 2+ (it fetches) for less effort — but then VERIFY every event time/number.
```

- Anchor it to a routine cue ("coffee → briefing"). It compresses gathering; you keep the thinking and still write your own bias + if–then plan.
- Avoid letting it replace your analysis, trusting unverified data, or asking it what to trade. Keep the briefing tight enough that you'll actually read it daily.

CHAPTER 4

The Trade-Plan Gatekeeper Agent

This is the agent that protects you from yourself. Most retail accounts don't die from bad analysis — they die from **broken discipline**: the impulsive entry, the oversized position, the revenge trade, the setup that skipped half the checklist. The **Trade-Plan Gatekeeper** runs your pre-trade checklist (**Vol 10, Ch. 4** / **Vol 4**) against *every* proposed trade — tirelessly, identically, immune to the emotions that make you cut corners.

> 🔑 **The Gatekeeper enforces rules you already wrote — it doesn't invent them.** It is not deciding whether the trade is "good." It is checking whether *this trade obeys the rules you set when you were calm*. That distinction is everything: it keeps the agent inside the co-pilot role (Ch. 1) and keeps your discipline intact when adrenaline is high.

4.1 What the Gatekeeper does

You describe a setup you're considering; it returns a structured **GO / NO-GO** verdict against your checklist:

```
┌─ TRADE CHECK — EUR/USD long ──────────────────────────────────────┐
│ □ Aligns with my macro/daily BIAS?          ✅ (bias: long per briefing)│
│ □ At a meaningful LEVEL?                     ✅ (H4 support 1.0820)      │
│ □ Real CONFLUENCE (≥2 factors)?             ✅ (support + trendline)    │
│ □ Valid TRIGGER present?                     ⚠️ ("forming," not closed) │
│ □ Stop at invalidation, reward ≥ 2R?        ✅ (SL 1.0795, TP 1.0890)  │
│ □ SIZED so stop = ≤ 1% risk?                ❌ (your size = 1.8% risk)  │
│ □ No high-impact event imminent?            ✅ (next event in 4h)       │
│ □ Calm — no FOMO / revenge / overtrading?   ⚠️ (3rd trade in 20 min)   │
│                                                                        │
│ VERDICT: ❌ NO-GO — 1 hard fail (size > 1% risk) + 2 cautions.         │
│ FIX: reduce size to ≤1% risk; wait for trigger candle to CLOSE; check │
│ whether this is overtrading. Re-run when addressed.                    │
└────────────────────────────────────────────────────────────────────────┘
```

The rule is simple and strict: **any hard box unchecked** → **NO-GO**. The agent's job is to make that rule un-skippable.

4.2 Why an AI Gatekeeper beats a paper checklist

You already have a checklist (Vol 4/10). Why route it through an agent?

```
A PAPER CHECKLIST                     THE GATEKEEPER AGENT
  • Easy to rush or "eyeball"           • Forces you to articulate each box in words
  • Skipped under emotion               • Runs identically every time, unemotionally
  • Silent on position size             • Computes the risk % from your numbers
  • Doesn't push back                   • Names the specific fail and the fix
  • No record                           • The check itself becomes a journal note (Ch. 5)
```

The act of *typing out* the setup to the Gatekeeper is itself a discipline circuit-breaker — it inserts a deliberate pause between impulse and execution, which is exactly where revenge and FOMO trades die.

> ⚠️ **The Gatekeeper checks rules, not market wisdom.** It cannot know if a setup will work — nothing can (Ch. 1). A "GO" verdict means *"this trade obeys your rules,"* not *"this trade will win."* You can follow every rule and still lose the trade; that's normal and fine (Vol 4 — it's about the process, not any single outcome). Never read "GO" as a prediction.

4.3 The position-size & risk function

The most valuable concrete thing the Gatekeeper does is **math you must never get wrong**: position sizing to a fixed risk (Vol 1, Ch. 5 / Vol 4).

```
You give:  account size, risk % (e.g. 1%), entry, stop, pair.
It returns: max risk in $, stop distance in pips, and the position size (lots)
            that keeps the loss at your stop ≤ your risk %.
It flags:  if the size you INTENDED exceeds your risk rule → hard NO-GO.
```

This single function prevents the most common account-killer: an oversized position. Still, **sanity-check the arithmetic** — verify against your broker's calculator the first several times (Ch. 1: verify everything).

4.4 The core prompt (starter)

A version usable in any chat app — full version in the **Prompt Pack**:

```
ROLE: You are my trade-plan gatekeeper. You check a proposed trade against MY fixed
rules and return GO / NO-GO. You do NOT predict outcomes or judge if it will "work."

MY RULES (from my Vol 7 plan):
  • Bias must align with my stated daily/macro bias
  • Entry at a meaningful level with ≥2 confluence factors and a valid, CLOSED trigger
  • Stop at invalidation; reward ≥ 2R
  • Risk ≤ 1% of account per trade; account = [$X]
  • No entries within [30] min of a high-impact event
  • Must be calm — not revenge/FOMO/overtrading (>[N] trades or after a big loss)

WHEN I DESCRIBE A TRADE:
  1. Check each rule; mark ✅ / ⚠️ / ❌ with a one-line reason.
  2. Compute the position size for ≤1% risk from my entry/stop/account.
  3. VERDICT: any hard ❌ → NO-GO. List the specific fixes.
  4. Be strict and unemotional. Never soften a hard fail.

DON'T: predict if it'll win, encourage a trade that breaks a rule, or invent rules.
```

4.5 Using it in the session (without slowing to a crawl)

```
SESSION FLOW (Vol 10, Ch. 4):
  Setup appears → describe it to the Gatekeeper → read verdict →
    GO  → execute exactly as planned (entry, stop, size, target)
    NO-GO → do NOT trade; note the fix; wait for a clean setup
```

For scalpers where seconds matter, you won't run it on every entry live — instead use it to **pre-validate your setups before the session** and as a **circuit-breaker** when you feel the urge to deviate. For day/swing trading, running it per trade is realistic and powerful.

> 🔑 **The Gatekeeper's highest value is at your weakest moment.** When you're tilted after a loss and itching for revenge, that's exactly when a calm, rule-bound second voice saves your account. Build the habit of running it *especially* when you least want to — that reluctance is the signal you need it most.

4.6 Chapter summary

```
• The Gatekeeper runs your pre-trade CHECKLIST (Vol 4/10) on every proposed trade and
  returns a strict GO / NO-GO. Any hard box unchecked → NO-GO. It enforces rules you
  wrote when calm; it does not invent them or predict outcomes.
• It beats a paper checklist: unemotional, identical every time, computes risk %, names
  the fail + fix, and the act of typing the setup is a discipline circuit-breaker.
• Its key function is POSITION SIZING to ≤1% risk — preventing the #1 account-killer.
  Verify the math against your broker's calculator early on.
• "GO" means "obeys your rules," NOT "will win." You can follow every rule and lose —
  that's normal; it's about process, not any single outcome.
• Run it per-trade for day/swing; for scalping, pre-validate setups and use it as a
  circuit-breaker. Its value is highest exactly when you're tilted and least want it.
```

CHAPTER 5

The Journal & Analytics Agent

If you build only one agent, build this one. The journal is the engine of improvement (**Vol 4, Ch. 6**) — but it's also the discipline most traders abandon, because logging trades and crunching stats is tedious. The **Journal & Analytics Agent** removes the tedium: you dump the raw details of a trade, it produces a clean structured log entry, and on demand it computes the performance metrics that tell you what's actually working.

> 🔑 **This is the highest-ROI agent because it converts scattered trades into compounding lessons.** A trade you don't journal is a tuition payment you got nothing for. The agent makes journaling so frictionless that you actually do it every time — and that consistent record is what turns random outcomes into a clear, improvable edge over time.

5.1 What it does — two jobs

```
JOB 1 — LOGGING (every trade):  raw notes/screenshot → clean, structured journal entry
JOB 2 — ANALYTICS (on demand):  your logged trades → KPIs, segmentation, patterns
```

Job 1 — structured logging

You give it the messy reality; it returns a consistent entry in your Vol 4 format:

```
INPUT (you, fast & messy):
  "long eurusd 1.0820, stop 1.0795, target 1.0890, 0.5 lot, london open,
   support bounce + trendline, exited 1.0870 took partial, felt a bit early"

OUTPUT (structured entry):
  Date/Time · Pair: EUR/USD · Direction: Long · Session: London open
  Setup: support bounce + trendline confluence (Vol 3)
  Entry 1.0820 / Stop 1.0795 / Target 1.0890 / Size 0.5 lot
  Risk: 25 pips · Result: +50 pips (partial) · R-multiple: +2.0R
  Plan followed? ✅ entry & stop per plan · ⚠️ exited early vs target
  Process grade: B+ · Emotion: slight impatience
  Lesson tag: #early-exit #patience
```

Job 2 — analytics

On request it turns your logged trades into the numbers that matter (Vol 4, Ch. 7):

```
PERFORMANCE (last 30 trades):
  Win rate 47% · Avg win +2.3R · Avg loss −1.0R · Expectancy +0.55R/trade
  Profit factor 1.9 · Max drawdown −6.2R · % plan-followed 82%
SEGMENTATION:
  Best setup: trendline-confluence (+0.9R/trade) · Worst: counter-trend (−0.4R/trade)
  Best session: London · Worst: NY-close · Most costly tag: #revenge (−4.1R total)
```

> 🔑 **Segmentation is where the gold is.** The single most valuable thing analytics reveals is *which* setups, sessions, and behaviors make or lose you money. Cut your worst-performing setup and stop trading your worst

> session, and your equity curve often transforms — without learning anything new. The agent surfaces this from your own data in seconds.

5.2 Why route journaling through an agent

```
MANUAL JOURNALING                    THE JOURNAL AGENT
  • Tedious → skipped                  • Fast: dump messy notes, get clean entry
  • Inconsistent format                • Identical structure every time (comparable)
  • Stats need a spreadsheet           • KPIs on demand in plain language
  • Patterns hard to see               • Segments by setup/session/emotion instantly
  • Honest grading is uncomfortable    • A neutral party grades PROCESS, not ego
```

The agent also grades on **process, not profit** (Vol 4/10) — the right standard. A losing trade perfectly executed is a *good* trade; a winning trade that broke your rules is a *bad* one. The agent, having no ego, scores it honestly if you tell it to.

5.3 The data question — where do trades live?

```
LEVEL 0–1:  Paste trades into the chat; ask for analytics on a batch you paste.
            Simple; keep your own running log file (Markdown/CSV) alongside.
LEVEL 2–3:  The agent reads/writes a structured file (CSV/Markdown) you keep —
            it appends each entry and reads the whole file for analytics.
            Best balance of power and ownership.
ADVANCED:   A script pulls your broker's trade history (export/API), and the agent
            analyzes it. Most automated; verify the import maps correctly.
```

> ⚠️ **Own your data and keep it clean.** Your journal is your most valuable trading asset — keep it in a portable format (CSV/Markdown) that you control, not locked inside a chat you'll lose. And never paste broker logins or account numbers; the trade details are enough (Ch. 8). Verify the agent's computed KPIs against a manual spot-check periodically — analytics on wrong data is worse than no analytics.

5.4 The core prompt (starter)

Full version in the **Prompt Pack**:

```
ROLE: You are my trading journal & analytics assistant. You structure my trades and
compute my performance honestly. You grade PROCESS, not profit.

LOGGING: When I paste raw trade notes, output a clean entry in MY format:
  Date/Time · Pair · Direction · Session · Setup · Entry/Stop/Target/Size ·
  Risk(pips) · Result · R-multiple · Plan-followed? · Process grade · Emotion · Lesson tag(s)

ANALYTICS: When I ask, compute from the trades provided:
  win rate, avg win/loss (R), expectancy (R/trade), profit factor, max drawdown,
  % plan-followed. Then SEGMENT by setup, session, and emotion/tag, and name my
  single best and worst categories with the numbers.

RULES: grade on whether I followed my plan, not whether I profited. Be honest and
neutral. Only use the trades I provide. Flag anything that looks like a data error.
```

5.5 Plugging it into the routine

```
POST-MARKET (Vol 10, Ch. 5):   each session close → dump the day's trades → log them
WEEKLY (Vol 10, Ch. 7):        ask for analytics on the week → feed the Review agent (Ch. 7)
MONTHLY:                        full-sample KPIs + segmentation → strategy decisions (Vol 4)
```

The Journal agent is the **data backbone** of your whole AI system: the Gatekeeper's checks, the day's trades, and your emotions all flow into it — and the Review agent (Ch. 7) reads it back to find your patterns. Build it first; everything else gets richer because of it.

5.6 Chapter summary

```
• The Journal & Analytics agent is the HIGHEST-ROI agent — build it first. Two jobs:
  (1) turn messy trade notes into clean structured entries, (2) compute KPIs on demand.
• It removes the tedium that makes traders abandon journaling, so you actually log every
  trade — and that consistent record is what compounds into an improvable edge.
• ANALYTICS: win rate, avg win/loss, expectancy, profit factor, drawdown, % plan-followed.
  SEGMENTATION (by setup/session/emotion) is where the gold is — cut your worst category.
• It grades on PROCESS, not profit — the right standard, scored honestly by an ego-free party.
• Keep your journal in a portable file YOU own (CSV/Markdown); never paste credentials;
  spot-check the computed KPIs against your data.
• It's the data backbone: it feeds the Review agent (Ch. 7) and informs strategy (Vol 4).
```

CHAPTER 6

The Macro & News Research Agent

The macro world (**Vol 6**) is a firehose: central-bank statements, data releases, geopolitics, intermarket signals — far more than any solo trader can read, let alone synthesize. The **Macro & News Research Agent** is your tireless research analyst. It digests the flood and translates it into the only thing you need: *"what does this mean for my pairs?"* — while you keep the judgment about whether and how to act.

> 🔑 **The research agent turns information into understanding, not into trades.** It explains *what happened* and *why it matters* for your currencies (Vol 6/8) — central bank tone, data surprise, risk sentiment. It does **not** tell you to buy or sell. You combine its synthesis with your own analysis to form a view. Understanding in, decision yours.

6.1 What it does

```
INPUT:  a statement, headline, data release, or a research question
OUTPUT: a structured, plain-English brief:
  • WHAT happened (the facts, concisely)
  • SURPRISE vs expectations (the part that actually moves price — Vol 6, Ch. 4)
  • WHY it matters (transmission to rates / growth / risk sentiment)
  • PAIR IMPACT (which of MY pairs it touches, and the mechanism — not a direction call)
  • WATCH NEXT (related events/levels to monitor)
  • CONFIDENCE & CAVEATS (what's uncertain; what to verify)
```

Example uses:

```
• "Summarize today's FOMC statement vs the last one — what changed in tone?" (Vol 9, Ch. 3)
• "CPI came in 3.4% vs 3.2% expected — what's the mechanism into USD?" (Vol 6, Ch. 3–4)
• "Explain the current risk-on/off regime and what's driving it." (Vol 6, Ch. 5)
• "How does a rising oil price transmit to CAD and JPY?" (Vol 8)
• "Give me the bull and bear case for EUR/USD over the next month." (balanced, both sides)
```

That last one is a power-move: ask for **both sides** of a thesis. A good research agent argues bull *and* bear, exposing what you might be missing — the opposite of an oracle that just picks a direction.

6.2 The live-data requirement

This agent is only as good as its information — and macro is time-sensitive:

```
✗ A plain LLM (no tools) has a KNOWLEDGE CUTOFF — it does NOT know today's CPI print,
  this morning's central-bank decision, or current prices. Ask it for "today's news"
  and it may CONFIDENTLY INVENT it. This is the most dangerous misuse in the volume.

✓ Use one of these instead:
  • Paste the actual source text (statement, release, article) → it summarizes THAT
  • Give it a web-search / news tool (Level 2+) → it fetches, you verify the source
  • Use it for TIMELESS understanding ("how do rate differentials drive FX?") where the
    cutoff doesn't matter
```

⚠️ **Never ask a cutoff-bound model for current facts and trust the answer.** "What did the ECB do today?" to a tool-less LLM is an invitation to hallucinate a decision that never happened. Either feed it the real source or give it a verified live tool — and even then, confirm market-moving facts at the primary source (the central bank, the official statistics agency). Verify, always (Ch. 1).

6.3 Where it's genuinely powerful

Used correctly, this agent gives a solo trader something close to an analyst desk:

```
• DECODING central-bank language — parsing hawkish/dovish nuance & what CHANGED (Vol 9, Ch. 3)
• EXPECTATIONS framing — separating the number from the SURPRISE that moves price (Vol 6, Ch. 4)
• INTERMARKET joins — connecting yields, oil, gold, equities to your pairs (Vol 6, Ch. 6 / Vol 8)
• SPEED — a readable brief from a long document in seconds, so you're never blindsided
• EDUCATION — on-demand explanations that deepen your macro fluency over time (Vol 6)
• BALANCE — bull-vs-bear framing that counters your own confirmation bias
```

The goal is not to outsource your macro view — it's to **build your macro fluency faster** and stay current with less effort, so your *own* view is better informed.

6.4 The core prompt (starter)

Full version in the **Prompt Pack**:

```
ROLE: You are my macro & news research analyst. You explain what happened and what it
means for MY pairs. You never give buy/sell calls or predict price.

MY PAIRS & THEMES: [e.g. EUR/USD, GBP/USD, USD/JPY; current theme: rate-cut timing]

WHEN I PASTE A SOURCE OR ASK A QUESTION, respond in this format:
  • WHAT happened (concise facts)
  • SURPRISE vs expectations (the price-moving part)
  • WHY it matters (mechanism: rates / growth / risk sentiment)
  • PAIR IMPACT (which of my pairs, and HOW — mechanism, not a direction call)
  • WATCH NEXT · CONFIDENCE & CAVEATS (what to verify)
If I ask for a thesis, give BOTH the bull and bear case.

RULES: only use sources I provide or verified tool results — never invent current facts.
Flag your knowledge cutoff if asked about recent events without a source. Mark uncertainty.
```

6.5 Plugging it into the routine

```
PRE-MARKET (Vol 10, Ch. 3):   decode an overnight statement → feed your Briefing (Ch. 3)
OFF-HOURS (Vol 10, Ch. 6):    deeper study — understand a theme, build macro fluency (Vol 6)
EVENT-DRIVEN:                  a big release drops → paste it → get the "so what for my pairs"
WEEKLY (Vol 10, Ch. 7):       map next week's macro calendar & themes for your watchlist
```

🔑 **Research informs your bias; it never becomes your bias.** The agent hands you a clean, balanced synthesis. You still decide what it means for *your* positioning, weigh it against your technicals (Vol 3), and own the call. An analyst gives the manager a report; the manager decides. You are the manager.

6.6 Chapter summary

```
• The Research agent is your tireless analyst: it digests macro/news (Vol 6/8) into a
  structured "what happened → surprise → why → PAIR IMPACT → watch → caveats" brief.
• It explains and translates; it does NOT give buy/sell calls. Ask it for BOTH bull and
  bear cases to counter your own bias.
• LIVE-DATA RULE: a tool-less LLM has a knowledge cutoff and will INVENT current news.
  Either paste the real source, give it a verified live tool, or use it for timeless
  understanding. Always confirm market-moving facts at the primary source.
• It's powerful for decoding central-bank language, expectations framing, intermarket
  links, speed, and education — building your macro fluency faster.
• It feeds the Briefing agent and your study; but research INFORMS your bias, it never
  BECOMES it. You're the manager; the agent writes the report.
```

CHAPTER 7

The Review & Coaching Agent

Improvement comes from the review cycle (**Vol 10, Ch. 7 / Vol 7**) — but honest self-review is hard: you're too close, too biased, and too easily satisfied or discouraged. The **Review & Coaching Agent** is your neutral performance coach. It reads your journal data (Ch. 5), runs your weekly/monthly review, spots the patterns you can't see in yourself, and acts as a patient study partner — without ever flattering or sugar-coating.

> 🔑 **A coach holds up a mirror you can't hold up yourself.** You will rationalize your mistakes, forget your recurring errors, and over- or under-react to recent results. A neutral agent looking only at your data names the pattern plainly — "your last 8 losses were all counter-trend trades after a losing day" — which is the exact feedback that drives improvement and which you'd never give yourself.

7.1 What it does — three roles

```
ROLE 1 — REVIEWER:  runs your weekly/monthly review on your journal data (Vol 10, Ch. 7)
ROLE 2 — PATTERN-SPOTTER:  finds recurring behavioral/setup patterns across many trades
ROLE 3 — STUDY PARTNER:  quizzes you, reviews your reasoning, teaches from your own trades
```

Role 1 — the structured review

Fed your week's trades (from the Journal agent, Ch. 5), it produces your Vol 10 review:

```
WEEKLY REVIEW — week of [date]
  KPIs: 14 trades · Win 50% · Expectancy +0.4R · PF 1.7 · Max DD −3R · Plan-followed 79%
  BEST decision: held the EUR/USD winner to target despite the urge to cut
  WORST mistake: 2 revenge trades Thursday after the morning loss (−2.3R)
  RECURRING pattern: discipline drops after the first loss of the day
  ONE adjustment for next week: hard stop for the day after 2 losses (Vol 10, Ch. 4)
  WEEK AHEAD: key events [..] · focus: protect post-loss discipline
```

Role 3 — the study partner

In off-hours (Vol 10, Ch. 6) it keeps you sharp:

```
• "Quiz me on my own trading plan rules." (drills retention)
• "Here's a trade I took — critique my reasoning." (catches flawed logic)
• "Walk me through this setup as if teaching a beginner." (Feynman test of mastery)
• "What concept from Vol 6 am I weakest on? Test me on it."
```

7.2 Why it works — the neutrality advantage

```
SELF-REVIEW ALONE                    THE COACHING AGENT
  • Biased, ego-protective             • Neutral — reports the data, not your story
  • Forgets recurring errors           • Sees across ALL logged trades at once
  • Over-weights recent results        • Weighs the full sample (Vol 4)
```

```
  • Vague resolutions                   • Forces ONE specific, testable adjustment
  • Inconsistent cadence                • Same structured review every week/month
```

The agent's value is precisely that it has **no ego to protect and perfect recall of your data**. It will tell you the uncomfortable truth — that your "great week" was one lucky outlier, or that your worst habit is back — if you instruct it to be honest (and you should).

> ⚠️ **Demand honesty, and still keep perspective.** Tell the coach explicitly to be candid, not encouraging — a flattering coach is useless. But also remember it's working from limited data and small samples; treat its read as a *well-informed second opinion to weigh*, not a verdict. And don't let a tough review tip into discouragement — the point is the *one adjustment*, then back to the routine (Vol 10).

7.3 The "one adjustment" discipline

The single most important output of every review is **one specific change** — not ten vague intentions (Vol 10, Ch. 7):

```
✗ "Trade better, be more patient, manage risk, stop revenge trading, journal more..."
   → too many → nothing changes

✓ "After 2 losses in a day, I stop trading for the day."
   → one concrete, testable rule → measurable next week → real improvement
```

Instruct the agent to always end a review with exactly one prioritized, testable adjustment, and to **check next review whether you followed it**. That closes the improvement loop (Vol 4/10).

7.4 The core prompt (starter)

Full version in the **Prompt Pack**:

```
ROLE: You are my trading performance coach and study partner. You are neutral and
candid — never flattering. You work from MY journal data and MY plan.

WEEKLY/MONTHLY REVIEW: given my trades, produce:
  KPIs · best decision · worst mistake · ONE recurring pattern · ONE specific, testable
  adjustment for next period · what to focus on. Then check if I followed LAST period's
  adjustment. Grade me on PROCESS (plan-followed), not profit.

PATTERN-SPOTTING: across all trades I give you, surface recurring behavioral or setup
patterns — especially ones costing me money (e.g. revenge, overtrading, counter-trend).

STUDY PARTNER: when asked, quiz me on my plan/playbooks, critique my trade reasoning,
and test my understanding. Be a tough but constructive teacher.

RULES: be honest even when uncomfortable. Note small-sample caveats. End reviews with
exactly ONE adjustment. Never tell me to trade a specific direction.
```

7.5 Plugging it into the routine

```
WEEKLY (Vol 10, Ch. 7):   feed the week's journal → run the review → log the ONE adjustment
MONTHLY:                  larger-sample review → keep/cut/refine setups (Vol 4, Ch. 8)
OFF-HOURS (Vol 10, Ch. 6): study-partner mode — quiz & reasoning critique
QUARTERLY:                "am I progressing toward my goals?" (Vol 5/7) — big-picture check
```

> 🔑 **The review agent closes the loop that turns experience into skill.** Trade → journal (Ch. 5) → review (this agent) → one adjustment → trade again. That loop, run every week for years, is the entire mechanism of getting better (Vol 10, Ch. 8). The agent makes the loop frictionless and honest — which is exactly why most traders, who skip it, never improve.

7.6 Chapter summary

```
• The Review & Coaching agent is your neutral performance coach: it runs your weekly/
  monthly review on your journal data (Ch. 5), spots recurring patterns, and acts as a
  study partner (quizzes, reasoning critique, teaching from your own trades).
• Its edge is NEUTRALITY + perfect recall of your data — it names uncomfortable patterns
  you'd rationalize away. Instruct it to be candid, not encouraging.
• Every review ends with exactly ONE specific, testable adjustment — and the agent checks
  next time whether you followed it. That closes the improvement loop.
• Keep perspective: small samples, a second opinion to weigh, not a verdict. Don't let a
  tough review become discouragement — make the one change, return to the routine.
• Trade → journal → review → one adjustment → repeat. That loop, run for years, IS how
  you improve (Vol 10, Ch. 8). The agent makes it frictionless and honest.
```

CHAPTER 8

Building, Orchestrating & Safeguarding

You now know *what* agents to build and *why*. This chapter is the *how*: the practical craft of building reliable agents, connecting them into a system, and — most importantly — the safety rails that keep the whole thing from hurting you. This is the chapter that separates a useful AI support system from an expensive, dangerous mess.

> 🔑 **The safeguards are not optional add-ons — they are the product.** Anyone can paste a prompt into a chatbot. The skill is building agents that are *reliable* (consistent, verifiable), *safe* (they can't take a decision or leak your data), and *sustainable* (you actually keep using them). Treat this chapter's rules as the load-bearing walls of the whole volume.

8.1 The anatomy of a reliable prompt

Every agent's quality is decided by its instructions. A reliable agent prompt has six parts:

```
1. ROLE      – a clear identity & single job ("You are my trade-plan gatekeeper")
2. CONTEXT   – your rules, plan, pairs, style (from Vol 7) – what it needs to do the job
3. TASK      – exactly what to do, step by step
4. FORMAT    – the precise output structure (a template/table) – for consistency
5. RULES     – hard constraints, especially what NOT to do (no predictions, no invented facts)
6. EXAMPLE   – (optional but powerful) one example of a good input→output, so it matches your style
```

> 🔑 **The "DON'T" rules are the most important lines in any trading-agent prompt.** Always include the negative constraints explicitly: *don't predict price, don't give buy/sell calls, don't invent facts, don't soften a rule violation, flag uncertainty*. These guardrails are what keep the agent a co-pilot (Ch. 1). An agent prompt without them will eventually drift into giving you "advice" you shouldn't take.

8.2 Choosing your tools — the maturity ladder, practically

```
LEVEL 0  Saved prompts in any chat app. $0–20/mo. Start here for ALL agents.
LEVEL 1  "Custom assistants" with a fixed system prompt you reuse (most chat apps support this).
LEVEL 2  Add tools: web search, an economic-calendar feed, a price API, file reading.
LEVEL 3  Persistent data: the agent reads/writes your journal & plan files.
LEVEL 4  Orchestration: a script (e.g. Python) calls an LLM API, pulls data, chains agents,
         writes outputs. Use the latest capable models. This is a software project – only
         climb here if you'll maintain it.
```

> ⚠️ **Don't out-engineer your own discipline.** The most common failure is spending weeks building a Level-4 system and zero days actually trading the routine it supports. A Level-0 saved prompt you use every day beats a sophisticated pipeline you're perpetually "finishing." Build the simplest version that works, use it daily, and climb only when a real friction demands it.

8.3 The orchestration layer (advanced)

If you do climb to a connected system, the agents form a pipeline around your day:

```
        ┌──────────── your data store (journal CSV, plan, rules) ──────────┐
        │                                                                   │
  RESEARCH ──► BRIEFING ──► [YOU form bias & plan] ──► GATEKEEPER ──► [YOU trade]
   (Ch.6)       (Ch.3)              │                   (Ch.4)            │
        ▲                           │                                     ▼
        │                       JOURNAL  ◄────────────────────────────  (Ch.5)
        │                        (Ch.5)                                   │
        └────────── REVIEW (Ch.7) ◄── reads journal weekly/monthly ───────┘
```

Even fully orchestrated, note the two **[YOU]** nodes that no automation crosses: forming the bias and pulling the trigger. The pipeline moves *information*; you move the *decisions*. Keep those nodes human forever.

8.4 The safety rails (read this twice)

These are the non-negotiables. Violating any one of them can cost you real money:

```
1. NO AUTONOMOUS TRADING. No agent connects to your broker to place trades on its own
   judgment. The decision and the click stay human (Ch. 1, 2). Capability ≠ permission.
2. VERIFY EVERY FACT. Event times, prices, rates, calculations — confirm at the source
   before acting. Assume the agent CAN hallucinate, because it can (Ch. 1).
3. NEVER PASTE SECRETS. No broker passwords, API trading keys, account numbers, or PII.
   Trade details and analysis are enough. Assume anything you paste could be stored.
4. MIND YOUR DATA PRIVACY. Use providers/settings that don't train on your inputs where
   possible; keep your journal in a file YOU own. Your edge & data are yours.
5. GUARD AGAINST OVER-RELIANCE. Keep your own skills sharp; you must be able to trade
   WITHOUT the agents. They augment a competent trader; they can't create one.
6. THE PRIME DIRECTIVE OVERRIDES ALL. Agents support the process; you own every decision.
```

⚠️ **The single most dangerous thing you can build is an autonomous trading bot acting on AI "predictions" — especially on a funded/prop account (Vol 5), where it can breach rules and lose your funding in minutes while you sleep.** If you ever automate *execution*, it must be the mechanical execution of rules YOU defined and tested (Vol 4) — never the AI's discretionary "view." When in doubt, keep the human in the loop.

8.5 Common failure modes (and the fix)

```
FAILURE                               FIX
  Hallucinated facts                   Verify at source; feed real data; demand "flag uncertainty"
  Vague, generic output                Tighter ROLE + FORMAT + an example in the prompt
  Agent drifts into giving calls       Add explicit "DON'T predict / no buy-sell" rules
  Over-reliance / skill atrophy        Periodically trade & analyze WITHOUT the agents
  Outsourcing the decision             Re-read the prime directive (Ch. 1); keep [YOU] nodes human
  Over-engineering, under-trading      Drop to Level 0; ship the simple version; use it daily
  Privacy leak                         Never paste secrets; check provider data settings
  Acting on stale (cutoff) info        Give live tools or paste sources; never trust "today's news"
                                        from a tool-less model
```

8.6 Cost, sustainability & fit

```
• COST: Level 0–1 is a flat subscription; Level 4 adds API + data-feed usage costs —
  track them so the system earns its keep (it should save you time, not drain capital).
• SUSTAINABILITY: an agent system you maintain for years beats a clever one you abandon
  (the Vol 10, Ch. 8 lesson applies to your tools too). Keep it simple enough to sustain.
• FIT: adapt every agent to YOUR style, pairs, and plan (Vol 7). A generic agent gives
  generic help; a personalized one (fed your rules and history) gives sharp help.
• EVOLVE: refine your prompts on evidence — when an agent's output misleads or misses,
  fix the prompt. Your agents improve the same way you do: through review (Ch. 7).
```

8.7 Chapter summary

```
• A reliable agent prompt has six parts: ROLE, CONTEXT, TASK, FORMAT, RULES, (EXAMPLE).
  The "DON'T" rules (no predictions, no invented facts, flag uncertainty) are the most
  important lines — they keep the agent a co-pilot.
• Tool maturity ladder L0→L4: start at L0 (saved prompts), climb only when real friction
  demands it. Don't out-engineer your discipline — a daily L0 prompt beats an unfinished L4.
• Orchestration chains the agents around your day, but two [YOU] nodes stay human forever:
  forming the bias and pulling the trigger. Information flows; decisions don't.
• SAFETY RAILS (non-negotiable): no autonomous trading; verify every fact; never paste
  secrets; protect data privacy; guard against over-reliance; prime directive overrides all.
• Know the failure modes (hallucination, drift into calls, over-reliance, over-engineering)
  and their fixes. Mind cost, sustainability, and personalization; refine prompts on evidence.
```

APPENDIX

Agent Prompt Pack & Blueprints

Copy, paste, and personalize. These are complete, ready-to-use system prompts for every agent in this volume. **Fill in every `[bracketed]` field** with your own plan (Vol 7) — an agent is only as good as the context you give it (Ch. 2, 8). Every prompt already includes the safety rails from Ch. 8; **do not delete the "RULES / DON'T" sections** — they are what keep each agent a co-pilot, not an oracle (Ch. 1).

> 🔑 **Personalize before you use.** Paste your real pairs, session, risk limits, and checklist from your **Vol 7** plan. A generic prompt gives generic help; a prompt fed your actual rules and style gives sharp, usable help.

A. Your shared context block (paste into every agent)

Define this once and prepend it to any agent so they all "know" your trading:

```
MY TRADING CONTEXT (from my Vol 7 plan):
  • Trader style: [e.g. day/swing trader, London–NY session]
  • Pairs I trade: [e.g. EUR/USD, GBP/USD, USD/JPY]
  • Account size: [$X]    • Risk per trade: [≤1%]    • Daily loss limit: [e.g. 3%]
  • My edge/strategy in one line: [e.g. trend-pullback to confluence levels]
  • My pre-trade checklist: [bias · level · ≥2 confluence · closed trigger · ≥2R · ≤1% risk · no
imminent news · calm]
  • My non-negotiable rules: [e.g. no trading 30 min around high-impact news; stop for the day after 2
losses]
```

B. Pre-Market Briefing Agent (Ch. 3)

```
ROLE: You are my pre-market briefing analyst. Your job is to PREPARE me to trade by
organizing information. You never predict the market or tell me what to trade.

[paste MY TRADING CONTEXT]

INPUTS I'll paste each morning: today's economic calendar, overnight/Asian price action
notes, the current DXY / yields / VIX / gold read, and my marked key levels.

DO:
  1. Organize everything into this BRIEFING format:
     1) CALENDAR (today's high-impact events for my pairs + exact times)
     2) OVERNIGHT (what moved & why, briefly)
     3) RISK REGIME (risk-on/off read from the inputs)
     4) KEY LEVELS (my marked S/R per pair)
     5) THEMES (current macro narrative, 1–3 bullets)
     6) WATCH (pairs/events most worth attention today)
     7) PROMPTS (3–4 sharp questions to help ME form my own bias & day plan)
  2. Flag conflicts and notable risk events explicitly. Keep it tight and skimmable.

RULES / DON'T: use ONLY the data I provide — never invent events, numbers, or prices.
Do not predict direction or give buy/sell calls. End with the section-7 questions for me.
```

C. Trade-Plan Gatekeeper Agent (Ch. 4)

```
ROLE: You are my trade-plan gatekeeper. You check a proposed trade against MY fixed rules
and return GO / NO-GO. You do NOT predict whether it will work or judge it as "good."

[paste MY TRADING CONTEXT]

WHEN I DESCRIBE A TRADE (pair, direction, entry, stop, target, size, setup, my state):
  1. Check each item of my checklist; mark ✅ / ⚠️ / ❌ with a one-line reason each:
     □ aligns with my stated bias  □ at a meaningful level  □ ≥2 confluence factors
     □ valid CLOSED trigger  □ stop at invalidation & reward ≥2R  □ size = ≤1% risk
     □ no high-impact event imminent  □ calm (no FOMO/revenge/overtrading)
  2. Compute the position size that keeps the loss at my stop ≤1% of my account, from my
     entry/stop/account. Show the math (risk $, stop pips, lots).
  3. VERDICT: ANY hard ❌ → NO-GO. List the specific fix(es). Be strict; never soften a fail.

RULES / DON'T: a "GO" means "obeys my rules," NOT "will win" — say so if I treat it as a
prediction. Don't predict outcomes, don't encourage a rule-breaking trade, don't invent rules.
Tell me to verify the position-size math against my broker's calculator.
```

D. Journal & Analytics Agent (Ch. 5)

```
ROLE: You are my trading journal & analytics assistant. You structure my trades and
compute my performance honestly. You grade PROCESS (did I follow my plan?), not profit.

[paste MY TRADING CONTEXT]

LOGGING — when I paste raw trade notes, output a clean entry in this format:
  Date/Time · Pair · Direction · Session · Setup · Entry/Stop/Target/Size · Risk(pips) ·
  Result · R-multiple · Plan-followed? (✅/⚠️/❌ + note) · Process grade (A–F) · Emotion ·
  Lesson tag(s)

ANALYTICS — when I ask, compute from the trades provided:
  win rate · avg win (R) · avg loss (R) · expectancy (R/trade) · profit factor ·
  max drawdown (R) · % plan-followed. Then SEGMENT by setup, session, and emotion/tag,
  and name my single BEST and WORST category in each, with the numbers.

RULES / DON'T: grade on process, not profit. Be honest and neutral. Use ONLY the trades I
provide. Flag anything that looks like a data-entry error. Never give buy/sell calls.
```

E. Macro & News Research Agent (Ch. 6)

```
ROLE: You are my macro & news research analyst. You explain what happened and what it
means for MY pairs. You never give buy/sell calls or predict price.

[paste MY TRADING CONTEXT]   Current theme(s) I'm tracking: [e.g. rate-cut timing]

WHEN I PASTE A SOURCE (statement/release/article) OR ASK A QUESTION, respond in this format:
  • WHAT happened (concise facts)
  • SURPRISE vs expectations (the price-moving part)
  • WHY it matters (mechanism: rates / growth / risk sentiment)
  • PAIR IMPACT (which of my pairs, and HOW it transmits — mechanism, not a direction call)
  • WATCH NEXT (related events/levels)
  • CONFIDENCE & CAVEATS (what's uncertain; what I should verify at the source)
If I ask for a thesis/outlook, give BOTH the bull and the bear case.
```

```
RULES / DON'T: use ONLY sources I provide or verified tool results — NEVER invent current
facts. If I ask about recent events without giving a source, state that you may be past your
knowledge cutoff and ask me to paste the source. Mark uncertainty. No buy/sell calls.
```

F. Review & Coaching Agent (Ch. 7)

```
ROLE: You are my trading performance coach and study partner. You are neutral and CANDID —
never flattering. You work from MY journal data and MY plan.

[paste MY TRADING CONTEXT]

WEEKLY/MONTHLY REVIEW — given my trades for the period, produce:
  • KPIs (win% · expectancy · profit factor · max DD · % plan-followed)
  • Best decision · Worst mistake
  • ONE recurring pattern (especially a costly behavioral one)
  • ONE specific, testable adjustment for next period
  • What to focus on next period
  Then check whether I followed LAST period's adjustment. Grade on PROCESS, not profit.

PATTERN-SPOTTING — across all trades I give you, surface recurring behavioral or setup
patterns that cost me money (revenge, overtrading, counter-trend, early exits, etc.).

STUDY PARTNER — when asked: quiz me on my plan/playbooks, critique my trade reasoning,
and test my understanding. Be a tough but constructive teacher.

RULES / DON'T: be honest even when uncomfortable; note small-sample caveats; end every
review with exactly ONE adjustment. Never tell me to trade a specific direction.
```

G. Build & rollout checklist

```
□ Wrote my shared CONTEXT block (A) from my Vol 7 plan
□ Built TIER 1 first: Journal (D) → Briefing (B) → Gatekeeper (C)   [Ch. 2]
□ Each agent: kept the RULES/DON'T section intact   [Ch. 8]
□ Anchored each agent to a routine CUE (Vol 10): coffee→briefing, pre-entry→gatekeeper,
  close→journal, Sunday→review
□ Chose a tool level (start L0: saved prompts) and a private place for my journal file
□ Set my VERIFY habit: confirm event times, prices, rates, and math at the source
□ Safety rails on: no autonomous trading · no secrets pasted · own my data · stay skill-sharp
□ Added TIER 2 (Research, Review) once Tier 1 is a habit
□ Scheduled to REFINE prompts on evidence (monthly)   [Ch. 7, 8]
```

H. The whole system on one card

```
┌──────────────────────────────────────────────────────────────────────────┐
│  AI supports my PROCESS; I own every DECISION. I verify its facts; I never  │
│  trade on its prediction.                                                   │
│                                                                             │
│  BRIEFING prepares my day → I form bias & plan → GATEKEEPER vets each trade │
│  → I trade → JOURNAL logs & measures → REVIEW finds the one lesson →        │
│  RESEARCH keeps me informed throughout. Repeat daily, refine over years.     │
│                                                                             │
│  No bot trades for me. No fact goes unverified. No decision leaves my hands. │
└──────────────────────────────────────────────────────────────────────────┘
```